

13 Databases and Software Development

13.1 Entity Relationship Modelling

An **entity** is a category of object about which data is to be recorded.

- Each entity in a database system has **attributes**.
- Each entity needs an **entity identifier** which uniquely identifies the entity.
- The **primary key** is the entity identifier in a **relational database**.

An **entity description** is written using the format.

EntityName (PrimaryKey, Attribute1, Attribute2)

Two entities are said to be **related** if they are linked in some way.

The **degrees of relationship** between two entities can be

- One-to-one
- One-to-many
- Many-to-many

An **entity relationship diagram** is a diagrammatic way of representing the relationship between the entities in a database. Shows the

- Degree of relationship.
- Name of the relationship.

Relational Database

In a relational database, a separate **table** is created for each entity identified in the system.

- Where a relationship exists between entities, an extra field called a **foreign key** links the two tables.
- A foreign key is an attribute that creates a join between two tables - it is the attribute that is **common to both tables**.
- The primary key in one table is the foreign key in the table to which it is linked.

Many-to-many Relationships

Tables cannot be linked directly in a many-to-many relationship. Instead, create a **link table** with two foreign keys, each linking to one of the two tables. The two foreign keys also act as the primary key of the table.

A primary key which consists of more than one attribute is called a **composite primary key**.

13.2 Relational Databases and Normalisation

A **relational database** is a collection of tables which relationships are modelled by shared attributes.

Tables can be linked through the use of **common attributes**. This attribute must be a primary key of one of the tables, and is known as a **foreign key** in the second table.

Normalisation

Normalisation is a process used to come up with the **best possible design** for a relational database. Tables are organised in a way that

- No data is **unnecessarily duplicated**.
- Data is **consistent** throughout the database. This means anomalies will not arise when the data is inserted, amended or deleted.

Consistency should be an automatic consequence of not holding any duplicated data.

- The structure of each table is enough to allow you to enter **as many or as few items** as required.
- The structure should enable a user to make all kinds of **complex queries** relating data from different tables.

1. A table is in **first normal form** if it contains **no repeating attribute** or group of attributes.
2. A table is in **second normal form** if it is in first normal form and contains **no partial dependencies**.

A partial dependency is where one or more attributes **depends on only part of the primary key**, which can only occur if the primary key is a **composite key**.

3. A table is in **third normal form** if it is in second normal form and contains **no non-key dependencies**.

A non-key dependency is one where the value of an attribute is determined by the value of another attribute which is not part of the key.

Advantages of Normalisation

- Easier to **maintain and change**.

Data integrity is maintained since there is no unnecessary duplication of data, it will also be impossible to reference a non-existing record on another table.

- **Faster** sorting and searching, as normalisation produce **smaller tables with fewer fields**, searching is faster because less data is involved.

Holding data once saves storage space.

- A normalised data base with **correctly defined relationships** between tables will not allow records in a table on the ‘one’ side of a one-to-many relationship to be deleted accidentally.

13.3 Introduction to SQL

Structured query language (SQL) is a **declarative language** used for querying and updating tables in a relational database.

SELECT FROM WHERE

```
SELECT      -- List of fields to be displayed
FROM        -- List of tables where the data come from
WHERE       -- List of search criteria
```

ORDER BY

```
ORDER BY    Field1, Field2    ASC/DESC
```

13.4 Defining and Updating Tables Using SQL

CREATE TABLE

Create a new database table.

```
CREATE TABLE TableName (
    Field1    INTEGER NOT NULL PRIMARY KEY,
    Field2    VARCHAR(20) NOT NULL,
    Field3    DATE
)
```

Data types include

Data type	Description
CHAR(<i>n</i>)	Character string of fixed length <i>n</i>
VARCHAR(<i>n</i>)	Character string of variable length, max <i>n</i>
BOOLEAN	True or false
INTEGER	Integer
FLOAT	Floating point number
DATE	Day, month and year values
TIME	Hour, minute and second values
CURRENCY	Format numbers in currency used in your region

ALTER TABLE

Add, delete or modify columns in an existing table.

```
ALTER TABLE TableName
-- Operation, such as
ADD      FieldName  VARCHAR(10)          -- or
DROP     COLUMN     ColumnName          -- or
MODIFY COLUMN      ColumnName VARCHAR(30) NOT NULL
```

Defining Link Table

```
CREATE TABLE TableName (
    Field1      CHAR(4) NOT NULL,
    Field2      INTEGER NOT NULL,

    FOREIGN KEY Field1 REFERENCES Table1(Field1),
    FOREIGN KEY Field2 REFERENCES Table2(Field2),
    PRIMARY KEY (Field1, Field2)
)
```

INSERT INTO

```
INSERT INTO TableName(Field1, Field2)
VALUES (123, 456)
```

UPDATE

```
UPDATE TableName
SET Field1 = 123, Field2 = 456
WHERE Field1 = 987
```

DELETE

```
DELETE FROM TableName
WHERE Field1 = 123
```

Client-server Databases

Database Management System (DBMS) provides an option for client-server operation.

- DBMS **server software** runs on the network server.
- DBMS **client software** runs on individual workstations.

The server software **processes requests** for data searches and reports that originate from individual work stations running DBMS client software.

If the DBMS does not have client-server capability, the entire database would be **copied to the workstation** and software held on the workstation would search for the requested data.

- A large amount of time is being spent of transmitting irrelevant data.
- A longer search with a less powerful machine.

Advantages of client-server databases

- The **consistency** of the database is maintained because only one copy of the data is held.
- An expensive resource can be made **available to large number of users**.
- **Access rights** and security is managed and controlled centrally.
- **Backup and recovery** can be managed centrally.

Problems with Client-server Databases

Allowing multiple users to **simultaneously update a database table** may cause one of the updates to be lost.

1. When an item is updated, the entire record is **copied to to the user's local memory** at the work station.
2. When the record is saved, the record is rewritten to the server.

Record Locks

Record locking is a technique of **preventing simultaneous access to objects** in a database in order to prevent updates being lost or inconsistencies in the data arising.

- A record is locked whenever a user retrieves it for editing or updating.
- Anyone else attempting to retrieve the same record is **denied access** until the transaction is completed or cancelled.

If two users are attempting to update two records, a situation can arise in which **neither can proceed**, known as **deadlock**.

- **Serialisation** is a technique which ensures the **transactions do not overlap in time** and therefore cannot interfere with each other or lead to updates being lost. A transaction cannot start until the previous one has finished.

Can be implemented using **timestamp ordering**.

- **Timestamp ordering** - whenever a transaction starts, it is given a timestamp, so if two transactions affect the same object, the transaction with the **earlier timestamp is applied first**.

To ensure that transactions are not lost

1. Each object in the database has a **read timestamp and a write timestamp**, which are updated whenever an object in a database is read or written.
2. When a transaction starts, it reads the data from a record causing the **Read timestamp to be set**.
3. Before it writes the updated data back to the record, it will **check the read timestamp**.
4. If this is **not the same** as the value that was saved when this transaction started, it will know that another transaction is also taking place on the record.

The problem can be identified and avoided.

- **Commitment ordering** is a serialisation technique used to ensure that **transactions are not lost** when two or more users are simultaneously trying to access the same database object.
 - Transactions are ordered in terms of their **dependencies on each other** as well as the time they were initiated.
 - Deadlocks can be prevented by **blocking one request** until another is completed.